

Ship Image Classification using Deep Learning

Ahsan Khan

Introduction to Deep Learning

University of Colorado Boulder

Dataset Overview



The dataset consists of images of various ship types, categorized into five classes:

- Cargo (1)
- Military (2)
- Carrier (3)
- Cruise (4)
- Tankers (5)

The dataset includes 6,252 images for training and 2,680 images for testing.

Project Objective

To develop and compare different deep learning models for accurate ship classification

Selecting the best performing model for potential real-world applications.

Create a scalable solution for high-volume data processing

Methodology

- Data Loading and Initial Exploration
- Exploratory Data Analysis (EDA)
- Data Preprocessing
- Model Development
- Model Comparison
- Hyperparameter Tuning
- Comparison and Evaluation of Results
- Conclusion

Data Loading and Initial Exploration

```
# Load the train.csv file
train_df = pd.read_csv('train/train.csv')

# Display the first few rows and basic information about the dataset
print(train_df.head())
print("\nDataset Info:")
print(train_df.info())

# Display the class distribution
print("\nClass Distribution:")
print(train_df['category'].value_counts())

# Check for missing values
print("\nMissing Values:")
print(train_df.isnull().sum())
```

```
      image  category
0  2823080.jpg        1
1  2870024.jpg        1
2  2662125.jpg        2
3  2900420.jpg        3
4  2804883.jpg        2
```

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6252 entries, 0 to 6251
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- ---
0 image 6252 non-null object
1 category 6252 non-null int64
dtypes: int64(1), object(1)
memory usage: 97.8+ KB
None

Class Distribution:
category
1 2120
5 1217
2 1167
3 916
4 832
Name: count, dtype: int64

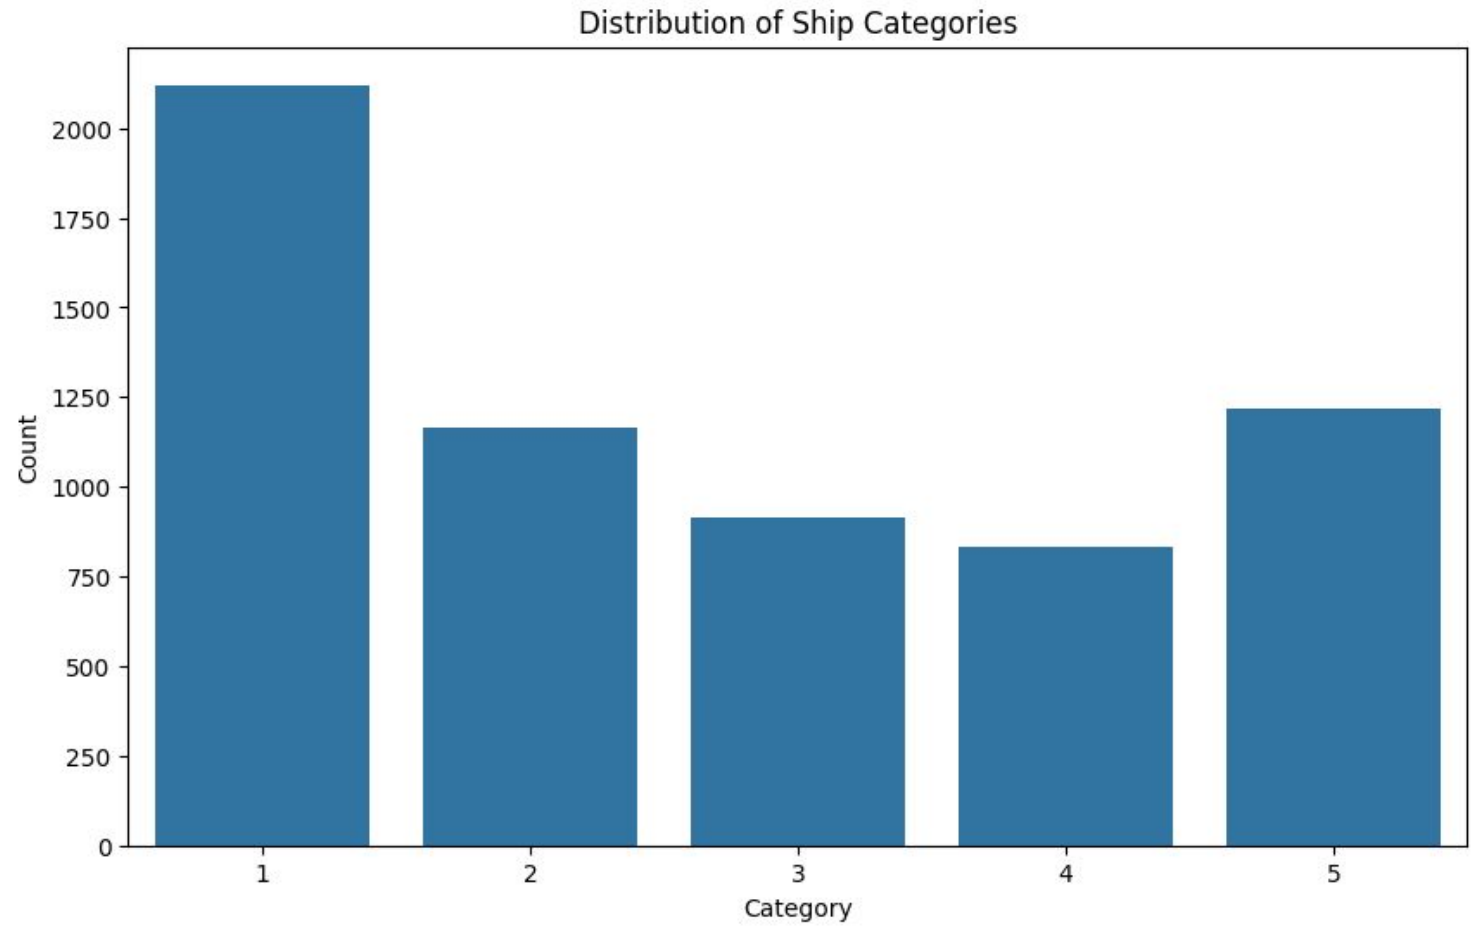
Missing Values:
image 0
category 0
dtype: int64

Data Cleaning

```
# Check for missing values  
print("\nMissing Values:")  
print(train_df.isnull().sum())
```

```
Missing Values:  
image      0  
category    0  
dtype: int64
```

Explorator y Data Analysis (EDA)



Explorator y Data Analysis (EDA)



Data Preprocessing

```
print("Number of training samples:", len(train_generator.filesnames))  
print("Number of validation samples:", len(val_generator.filesnames))  
print("Number of classes:", len(train_generator.class_indices))  
print("Class mapping:", train_generator.class_indices)
```

Found 5001 validated image filenames belonging to 5 classes.

Found 1251 validated image filenames belonging to 5 classes.

Category: 2



Category: 3



Category: 5



Category: 1



Category: 5





Number of training samples: 5001

Number of validation samples: 1251

Number of classes: 5

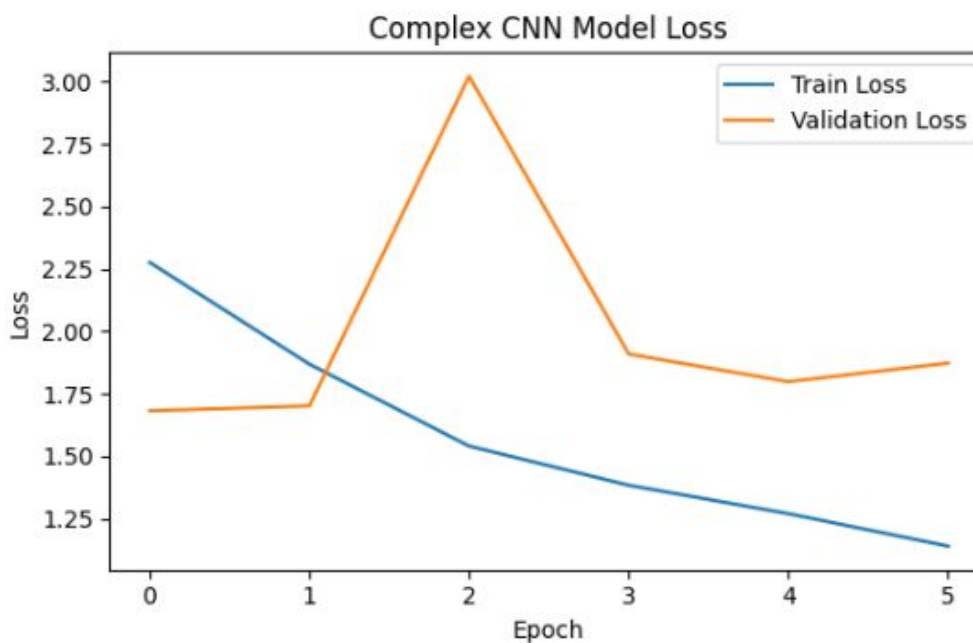

Class mapping: {'1': 0, '2': 1, '3': 2, '4': 3, '5': 4}

Simple CNN Model



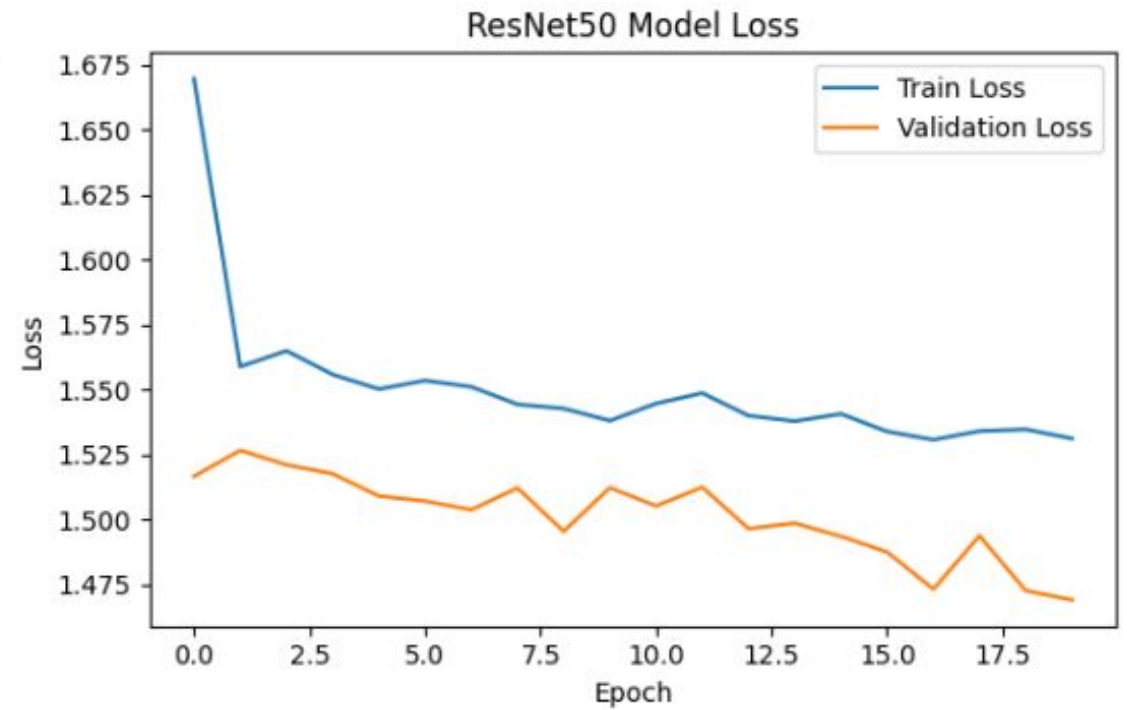
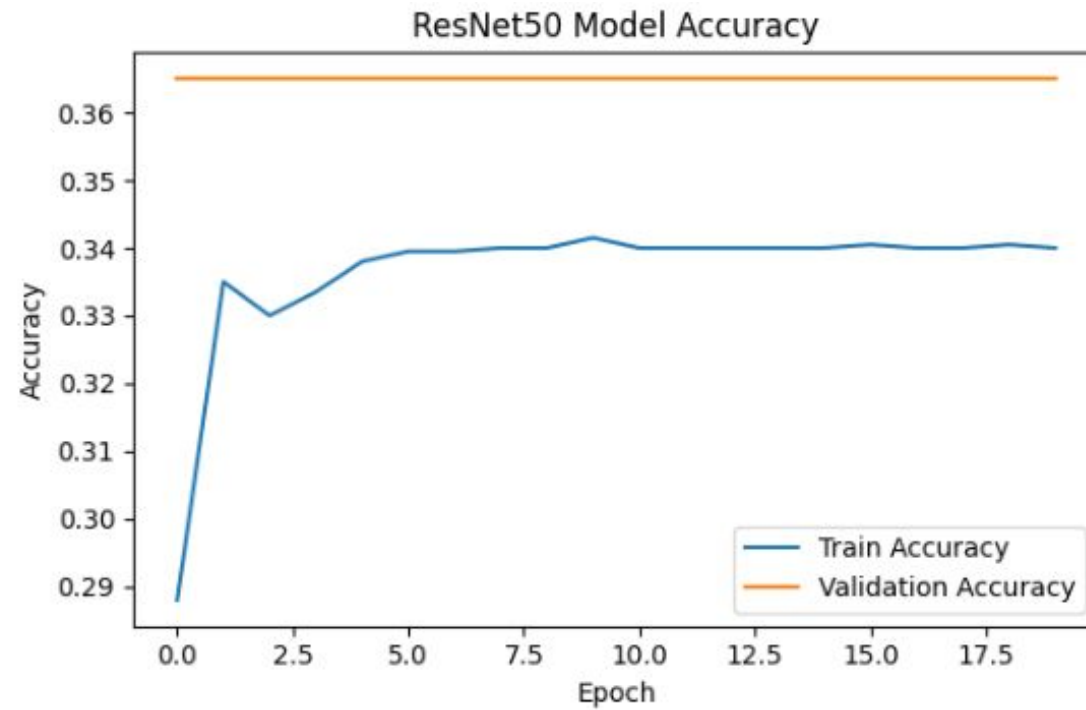
7/7 ————— 0s 34ms/step - accuracy: 0.4470 - loss: 1.5508
Test accuracy: 0.4600
Simple CNN - Final training accuracy: 0.5100
Simple CNN - Final validation accuracy: 0.4600
Simple CNN - Test accuracy: 0.4600

Complex CNN Model





10/10 — 3s 266ms/step - accuracy: 0.3596 - loss: 1.6528
Complex CNN Test accuracy: 0.3450

ResNet50 Model




13/13 — 25s 2s/step - accuracy: 0.3560 - loss: 1.4773
ResNet50 Test accuracy: 0.3650

MobileNetV2 Model



19/19 — 13s 690ms/step - accuracy: 0.8151 - loss: 0.4288
MobileNetV2 Test accuracy: 0.8150

Model Accuracy Comparison



Model Comparison Summary

	Model	Test Accuracy	Best Val Accuracy	Epochs Trained	\
3	MobileNetV2	0.815	0.826667	24	
0	Simple CNN	0.460	0.575000	20	
2	ResNet50	0.365	0.365000	20	
1	Complex CNN	0.345	0.345000	6	

	Training Time
3	276
0	190
2	190
1	15

Hyperparameter Tuning





```
Trial 5 Complete [00h 03m 37s]  
val_accuracy: 0.7549999952316284
```

```
Best val_accuracy So Far: 0.7699999809265137  
Total elapsed time: 01h 38m 28s
```

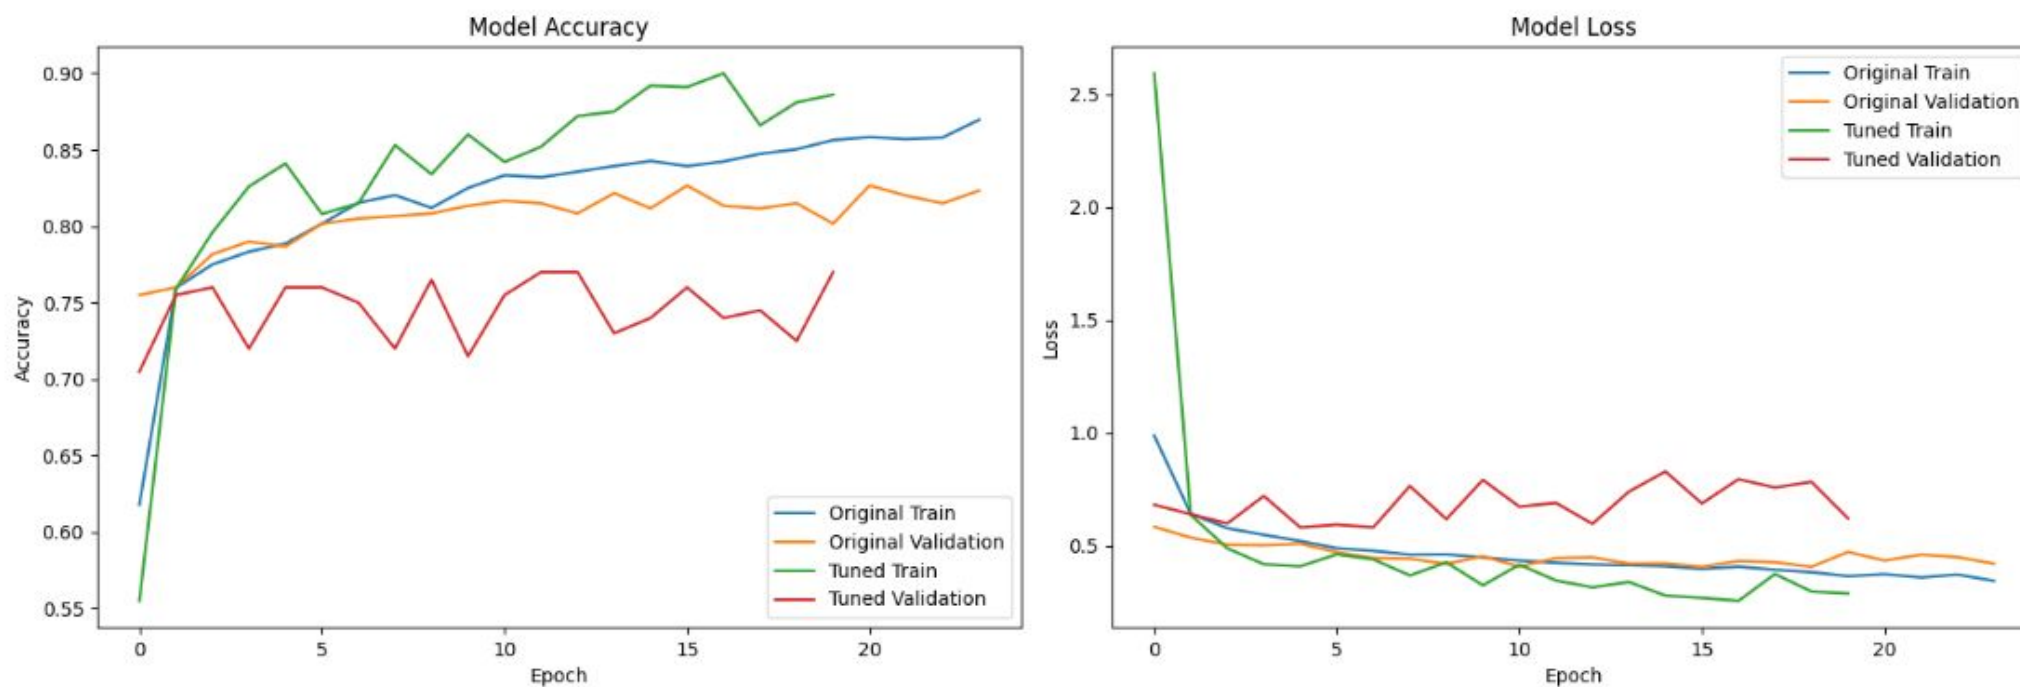
```
The hyperparameter search is complete. The optimal number of units in the dense layer is 480 and the optimal learning rate for the optimizer is 0.0089.  
The optimal dropout rate is 0.00.
```


Hyperparameter Tuning

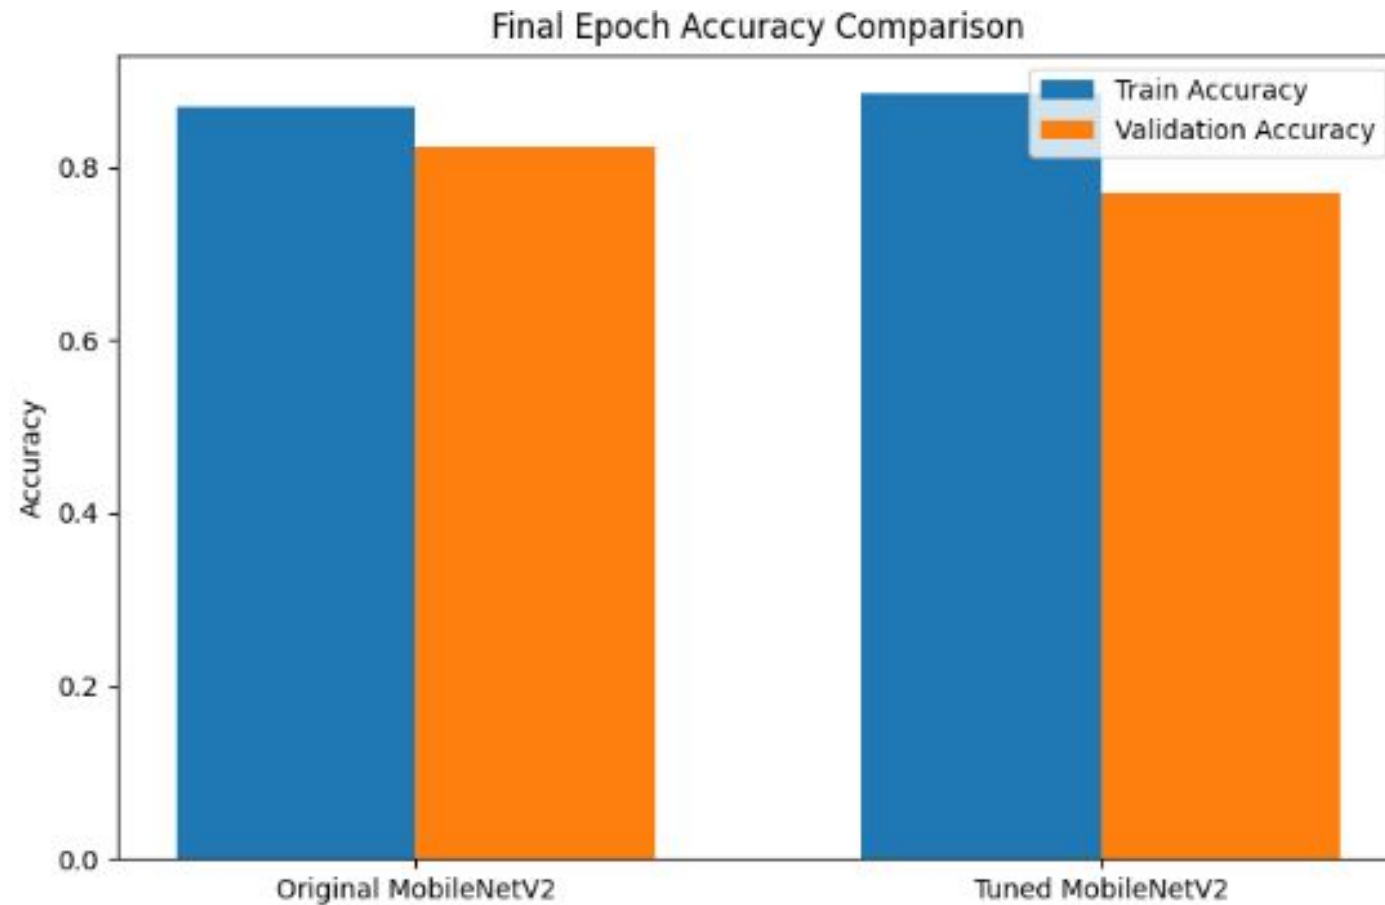
Validation accuracy: 0.7700



Hyperparameter Tuning Comparison



Hyperparameter Tuning Comparison



Conclusion

- The MobileNetV2 model demonstrated the best balance between performance and efficiency for this ship classification task.
- Transfer learning proved highly effective, significantly outperforming custom-built CNN architectures.
- The project highlighted the importance of model selection and the potential of lightweight architectures like MobileNetV2 for practical applications.